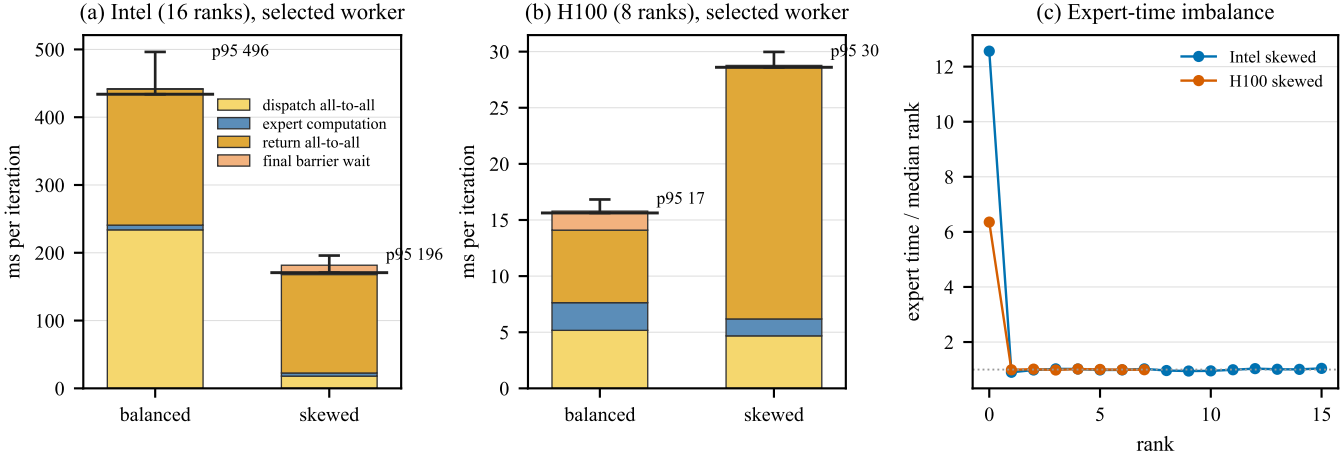
(a) Intel (16 ranks), selected worker
ms per iteration
p95 496
p95 196
dispatch all-to-all
expert computation
return all-to-all
final barrier wait
balanced
skewed
(b) H100 (8 ranks), selected worker
ms per iteration
p95 17
p95 30
balanced
skewed
(c) Expert-time imbalance
expert time / median rank
Intel skewed
H100 skewed
rank